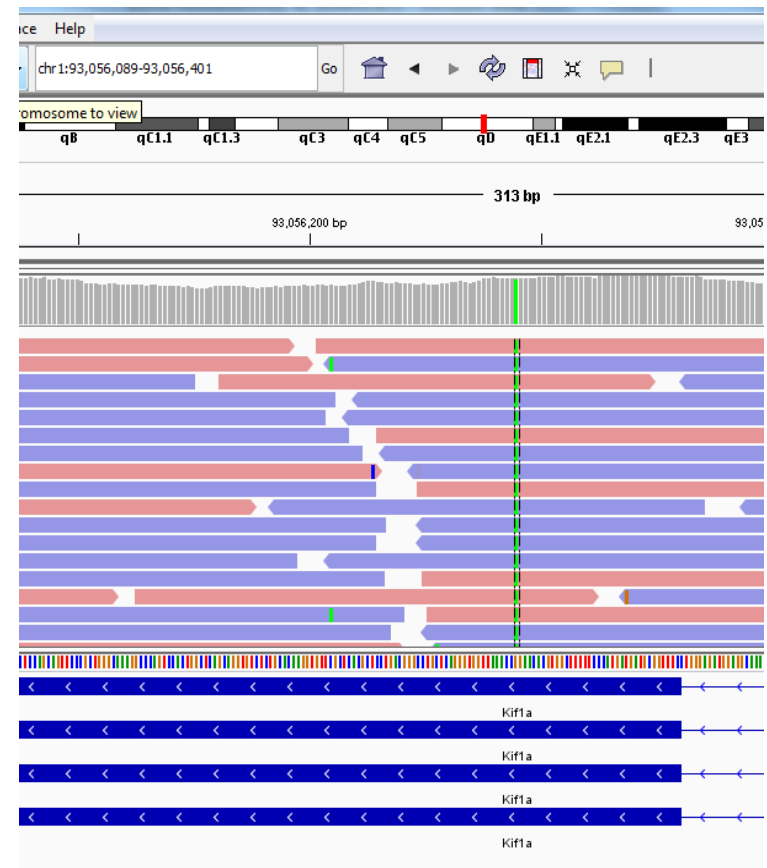




# High Throughput Sequence & Sequence Analysis: A General Review



# Sequencing - Methods

## Chain termination (Sanger sequencing)

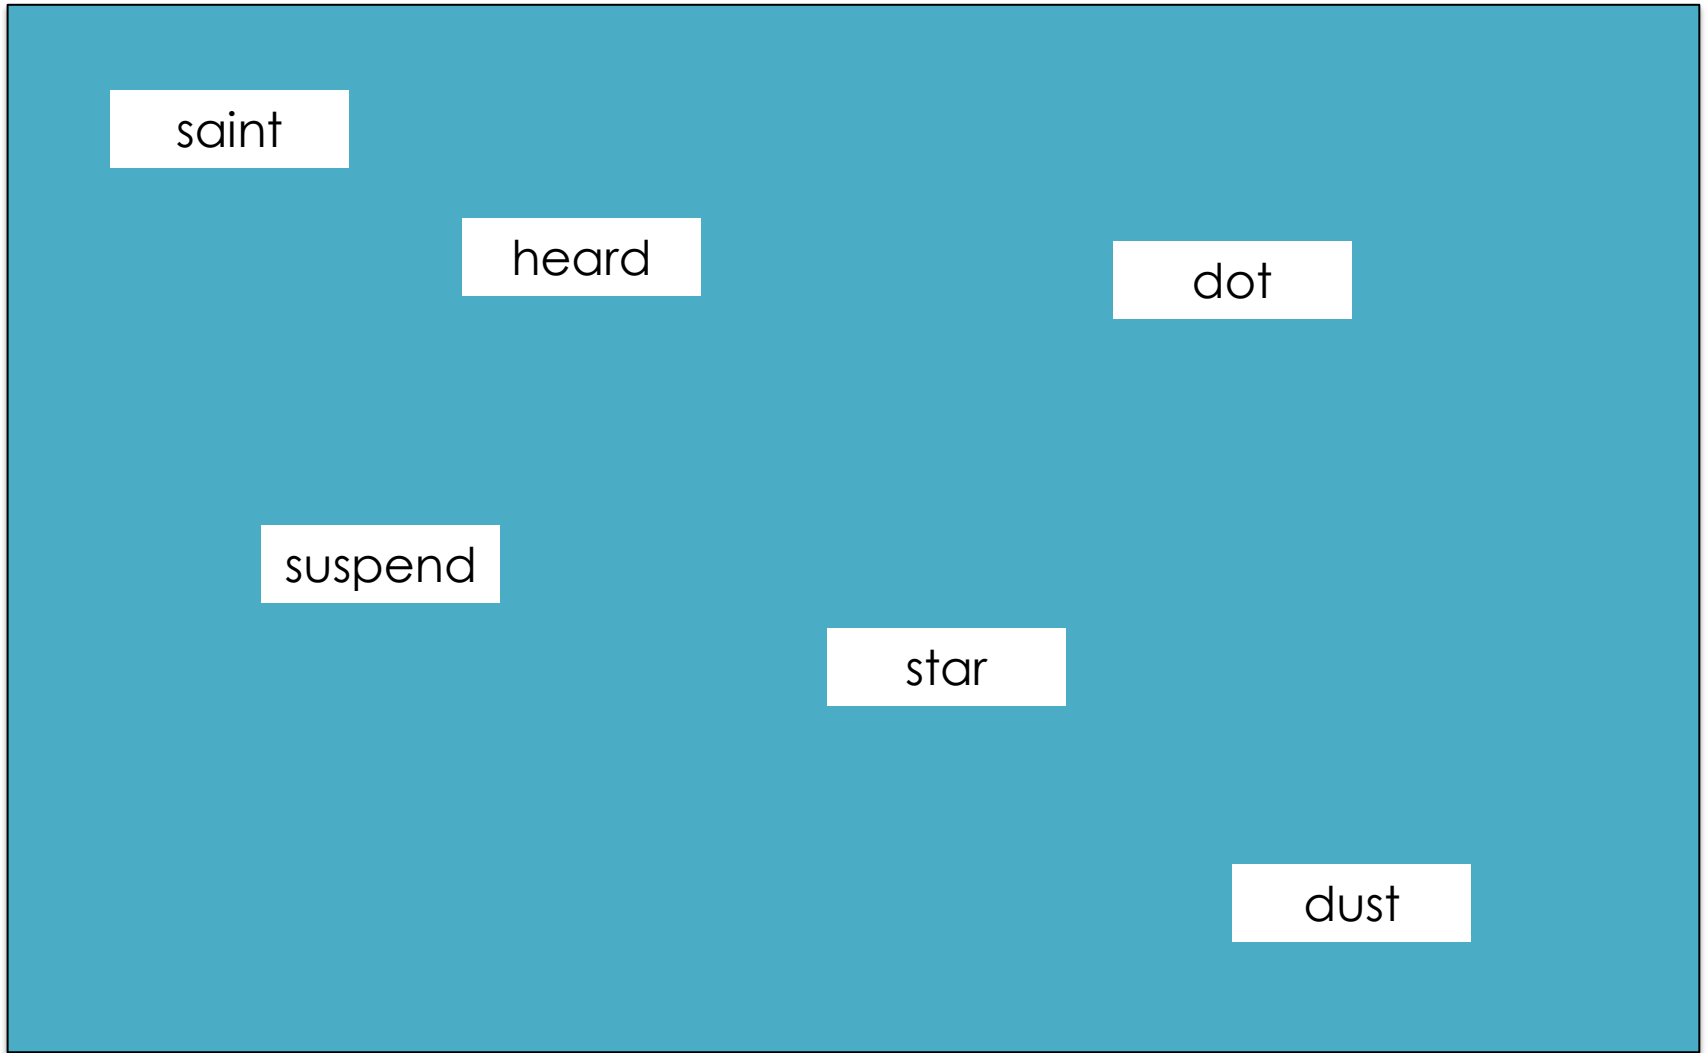


# Targeted sequencing limitations

	<i><b>Sanger</b></i>	<i><b>Pyro</b></i>
<i>Max. Length</i>	800-1000 bp	350-500 bp
<i>Error rate</i>	0.001% to 1%	> 1%


✧ **TARGETED!**

# How to build a genome by sequencing




✧ Where do the targeted sequences go in the genome?


# How does next-gen sequencing work?



# Terminology: libraries, lanes, and flowcells



Each reaction produces a unique **library** of DNA fragments for sequencing.




Each NGS machine processes a single **flowcell** containing several independent **lanes** during a single sequencing run


[http://www.illumina.com/technology/next-generation-sequencing/paired-end-sequencing\\_assay.html](http://www.illumina.com/technology/next-generation-sequencing/paired-end-sequencing_assay.html)

# High throughput sequencing

## Illumina sequencing – chain termination



# 2nd generation sequencing output formats



**Illumina**



**FASTQ (various flavours)**



**SoLID/ABI-Life**



**Color-space FASTA**



**Roche 454**



**SFF**



**Ion Torrent**



**SFF or FASTQ**



# Platforms have errors and artifacts



**Illumina**





**SoLID/ABI-Life**



**Roche 454**



**Ion Torrent**



**Removal of low quality bases**  
**Removal of adaptor sequences**  
**Platform specific artifacts (e.g homopolymers)**

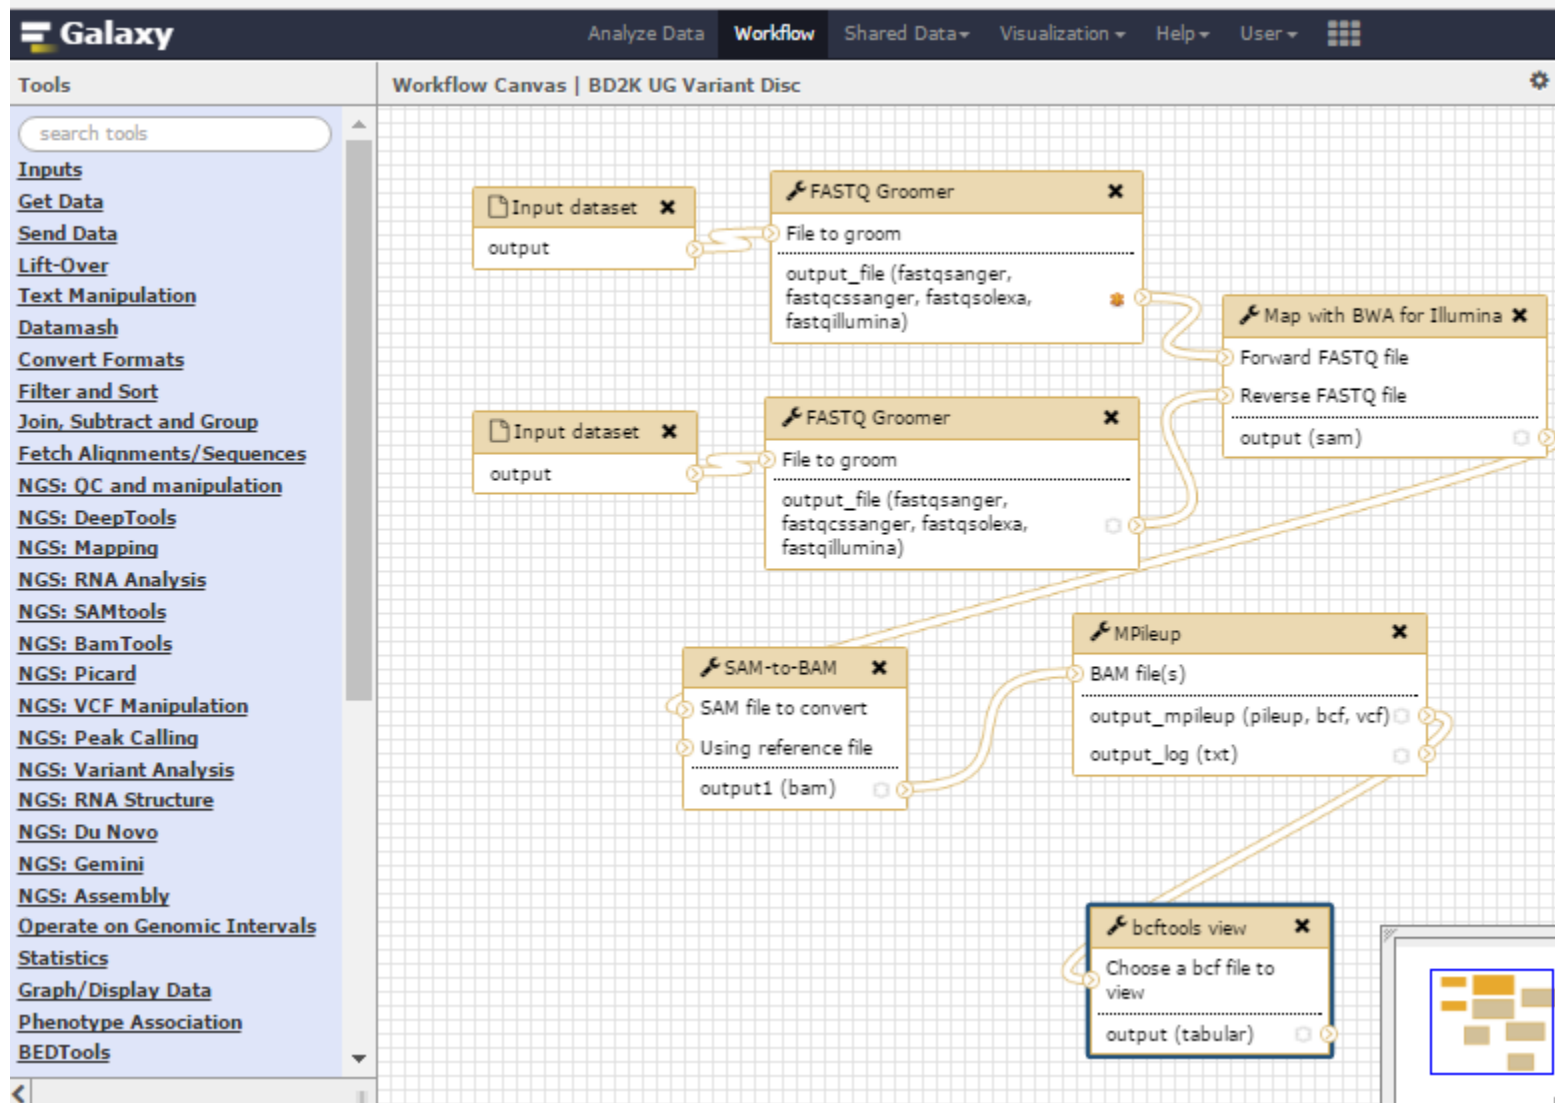
# Computational choices in genomic data analysis

- Download and run at the command-line
- Public web-based server
- Private instance of a web-based server

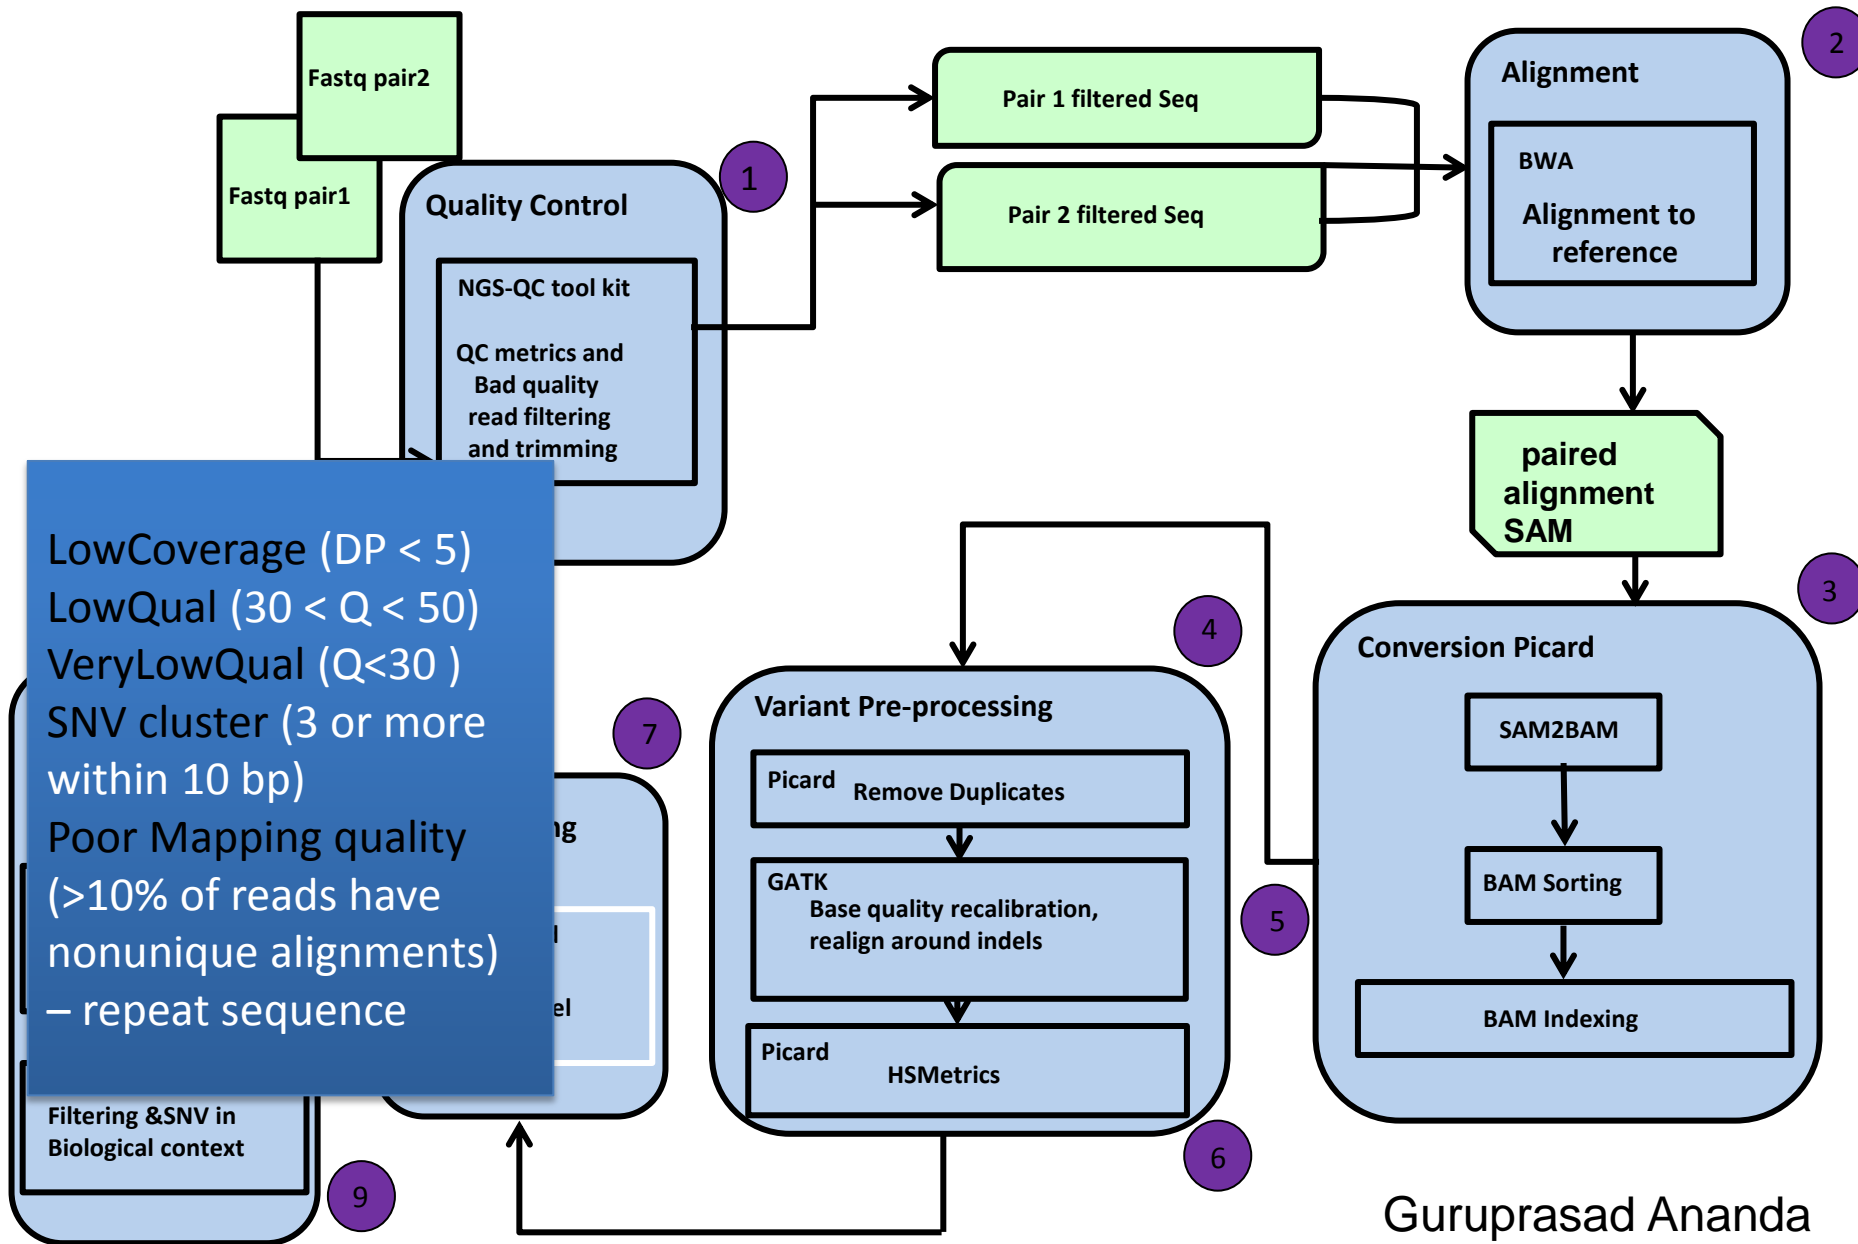
# Things to remember about server-based analysis

- You are using someone else's compute resources
  - Good- your investment is small, someone else
  - Bad- server load and wait times can be unpredictable
- The majority of the tools are “wrapped” presenting a reduced set of options/functions
  - Full functionality is accessed from the command-line
- File transfer times to and from the server will be dependent upon network capacity.

# Designing Data Analysis Workflows



# Mouse exome CIVET pipeline



# How to build a genome by sequencing

saint

heard

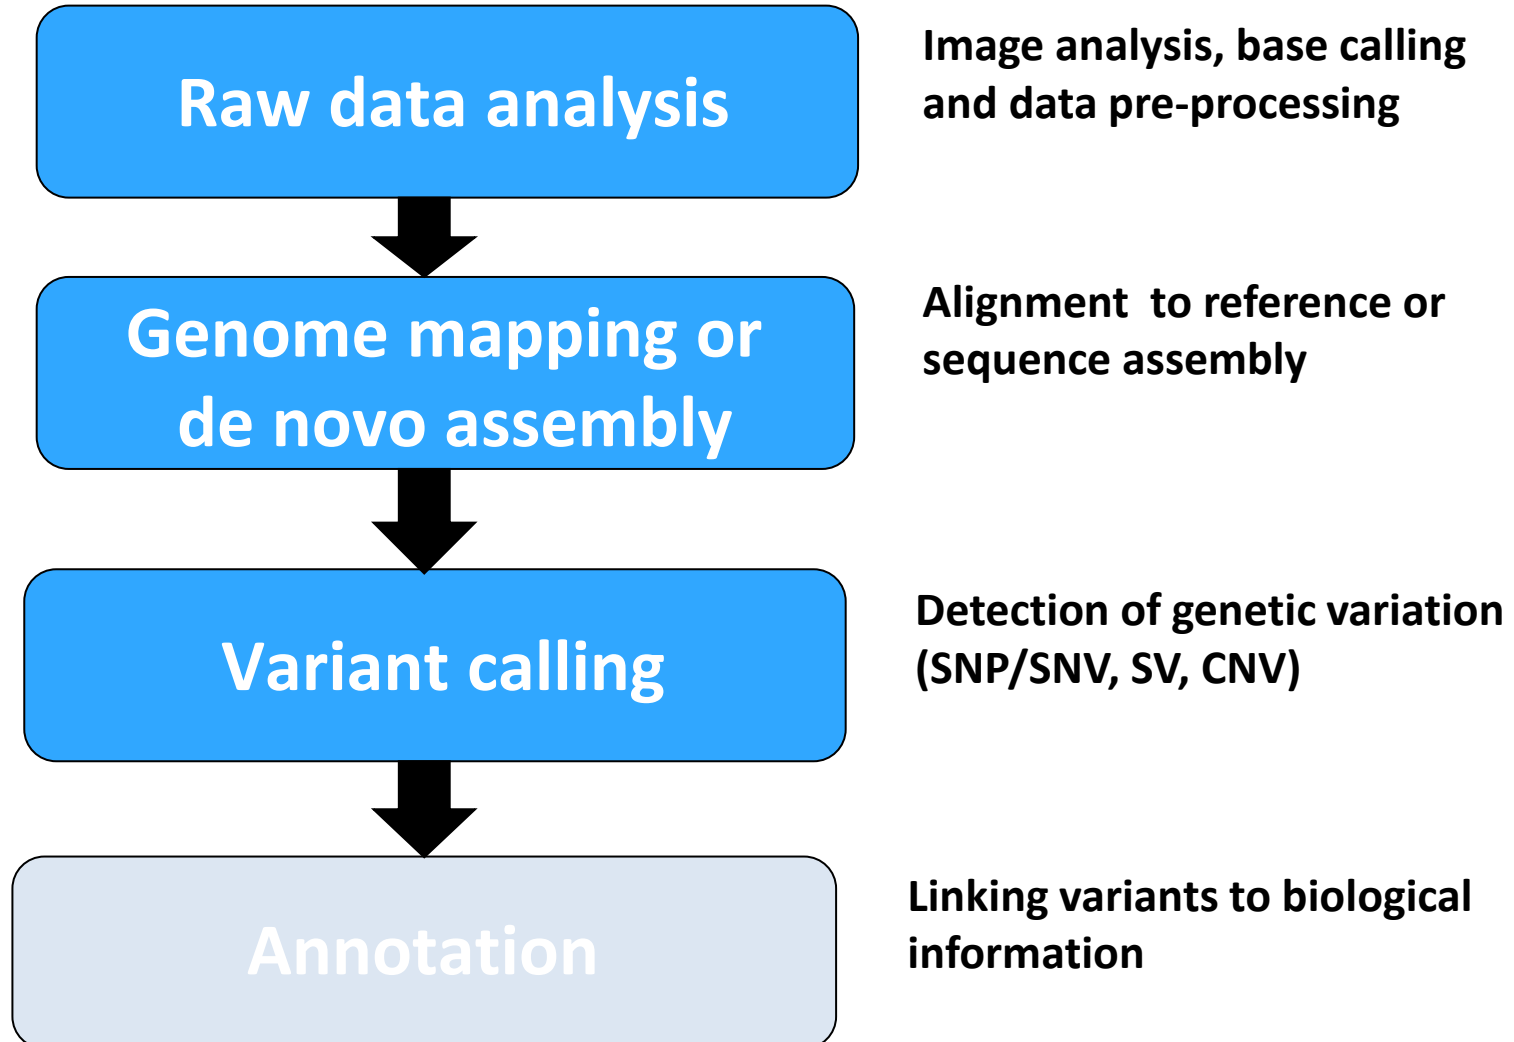
dot

suspend

star

dust

# Basic NGS data analysis workflow



# Fastq format

```
@HISEQ2000:128:D230MACXX:3:1101:1083:2161 1:Y:0:TAGCTT
CCATAGAAAGACTGGTTTGNNNANNNNNNNNNNNNNNNNNNNNNNNNNNNNNNNNNNNNNNNNNNN
+
; ; 5 = ; ? > ? > @ > < > ? ? > > @ ? ### 3 #####
```

- |                  |  |
|------------------|--|
| 1. @HISEQ2000    | the unique instrument name   |
| 2. 128:D230MACXX | the run id, the flowcell id  |
| 3. 3             | flowcell lane  |
| 4. 1101          | tile number within the flowcell lane   |
| 5. 1083:2161     | 'x'-coordinate of the cluster<br>within the tile:'y'- coordinate of the cluster<br>within the tile |
| 6. 1             | the member of a pair, 1 or 2<br>(paired-end or mate-pair reads only)                               |
| 7. Y             | Y if the read fails filter (read is bad),<br>N otherwise   |
| 8. 0             | 0 when none of the control bits are<br>on, otherwise it is an even number                          |
| 9. TAGCTT        | index sequence   |



# Quality value interpretation

Phred Quality Score	Probability of incorrect base call	Base call accuracy
10	1 in 10	90 %
20	1 in 100	99 %
30	1 in 1000	99.9 %
40	1 in 10000	99.99 %
50	1 in 100000	99.999 %

$$Q = -10 \log_{10} P$$

**Q = Phred Quality Scores**

**P = Base-calling error probabilities**

# FASTQC

Analyze Data   Workflow   Shared Data   Visualization   Help   User

Using 75.8 GB

Tools

search tools

[Get Data](#)  
[Send Data](#)  
[Text Manipulation](#)  
[Filter and Sort](#)  
[Join, Subtract and Group](#)  
[Convert Formats](#)  
[Extract Features](#)  
[Fetch Sequences](#)  
[Fetch Alignments](#)  
[Statistics](#)  
[Graph/Display Data](#)  
[Operate on Genomic Intervals](#)  
[Phenotype Association](#)  
[snpEff](#)  
[BEDtools](#)  
[Genome Diversity](#)  
[EMBOSS](#)  
[Regional Variation](#)  
[FASTA manipulation](#)  
[Evolution](#)  
[Multiple Alignments](#)

FastQC Read Quality reports (Galaxy Version 0.63)

Versions   Options

Short read data from your current history

1: MJFF2\_S1\_L001\_R1\_001.fastq

Contaminant list

Nothing selected

tab delimited file with 2 columns: name and sequence. For example: Illumina Small RNA RT Primer CAAGCAGAAGACGGCATACGA

Submodule and Limit specifying file

Nothing selected

a file that specifies which submodules are to be executed (default=all) and also specifies the thresholds for the each submodules warning parameter

Execute

**Purpose**

FastQC aims to provide a simple way to do some quality control checks on raw sequence data coming from high throughput sequencing pipelines. It provides a modular set of analyses which you can use to give a quick impression of whether your data has any problems of which you should be aware before doing any further analysis.

The main functions of FastQC are:

Import of data from BAM, SAM or FastQ files (any variant)  
Providing a quick overview to tell you in which areas there may be problems  
Summary graphs and tables to quickly assess your data  
Export of results to an HTML based permanent report  
Offline operation to allow automated generation of reports without running the interactive application

History

11: FastQC on data 2: Webpage

10: FastQC on data 1: RawData

9: FastQC on data 1: Webpage

312.7 KB  
format: html, database: ?

Run this job again


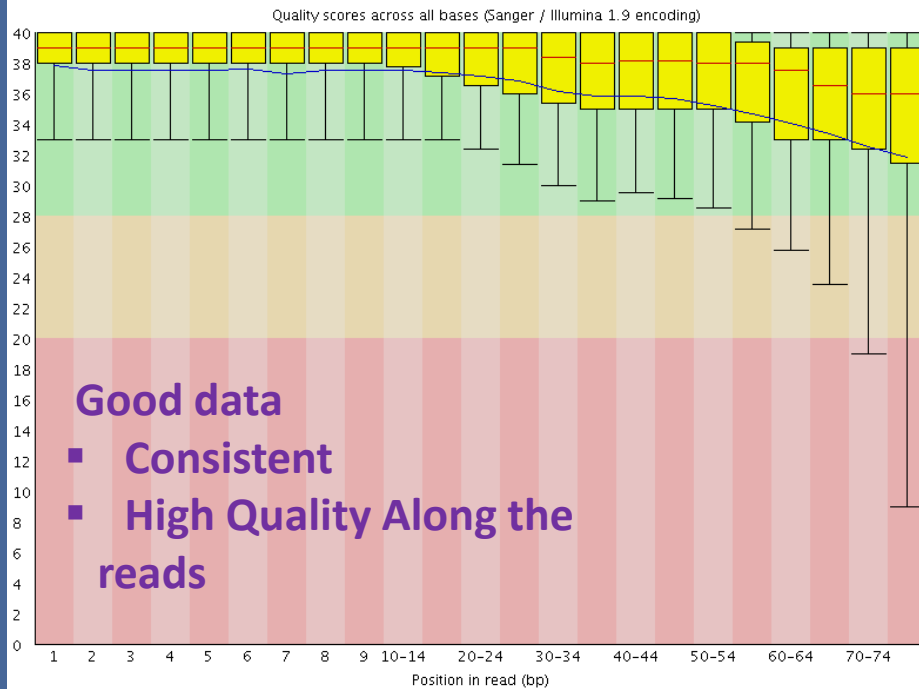
8: Line/Word/Character count on data 4

7: Line/Word/Character count on data 3

6: Line/Word/Character count on data 2

5: Line/Word/Character count on data 1

# Sequence quality per base position




## Bad data


- High Variance
- Quality Decrease with Length

- ❖ The central red line is the median value
- ❖ The yellow box represents the inter-quartile range (25-75%)
- ❖ The upper and lower whiskers represent the 10% and 90% points
- ❖ The blue line represents the mean quality

# Per sequence quality distribution



# Per sequence quality distribution



# Galaxy tool for filtering and trimming

**Quality Filter**

**Library to filter:**

**Quality cut-off value:**

**Percent of bases in sequence that must have quality equal to / higher than cut-off value:**

**FASTQ Quality Filter**

**Trim**


**Library to clip:**

**First base to keep:**

**Last base to keep:**

**FASTA/Q Trimmer**

# Nucleotide content per position



# Trimmomatic

**Trimmomatic flexible read trimming tool for Illumina NGS data (Galaxy Version 0.32.3)**

Options




**Paired end data?**

YesNo

**Input Type**

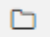
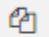

Pair of datasets

**Input FASTQ file (R1/first of pair)**



1: MJFF2\_S1\_L001\_R1\_001.fastq

**Input FASTQ file (R2/second of pair)**



2: MJFF2\_S1\_L001\_R2\_001.fastq

**Perform initial ILLUMINACLIP step?**

YesNo

Cut adapter and other illumina-specific sequences from the read

**Trimmomatic Operation**

1: Trimmomatic Operation

**Select Trimmomatic operation to perform**

Cut bases off the end of a read, if below a threshold quality (TRAILING)


**Minimum quality required to keep a base**

20

Bases at the end of the read with quality below the threshold will be removed



Instruments generate short reads that need to be mapped to the reference



# *Aligning (mapping) short sequencing reads*

1. Drop of water in the sea problem
2. Identical (or near identical) drops
3. Drops of polluted water

# Aligning (mapping) short sequencing reads

“Look again at that dot. That's here. That's home. That's us. On it everyone you love, everyone you know, everyone you ever heard of, every human being who ever was, lived out their lives. The aggregate of our joy and suffering, thousands of confident religions, ideologies, and economic doctrines, every hunter and forager, every hero and coward, every creator and destroyer of civilization, every king and peasant, every young couple in love, every mother and father, hopeful child, inventor and explorer, every teacher of morals, every corrupt politician, every "superstar," every "supreme leader," every saint and sinner in the history of our species lived there--on a mite of dust suspended in a sunbeam.”

— [Carl Sagan, \*Pale Blue Dot: A Vision of the Human Future in Space\*](#)

y m                      olo                      rin                      he                      ery

# Aligning (mapping) short sequencing reads

“Look again at that dot. That's here. That's home. That's us. On it everyone you love, everyone you know, everyone you ever heard of, every human being who ever was, lived out their lives. The aggregate of our joy and suffering, thousands of confident religions, ideologies, and economic doctrines, every hunter and forager, every hero and coward, every creator and destroyer of civilization, every king and peasant, every young couple in love, every mother and father, hopeful child, inventor and explorer, every teacher of morals, every corrupt politician, every "superstar," every "supreme leader," every saint and sinner in the history of our species lived there--on a mite of dust suspended in a sunbeam.”

— [Carl Sagan, \*Pale Blue Dot: A Vision of the Human Future in Space\*](#)

y m                      olo                      ron                      he                      ery

# DNA Sequence alignment

- sequence analysis → sequence alignment
- what:

```
aca--gacgcagtactttg-g-gc-caga-ac-cgt
      |||   |||  ||||  |  ||  ||||  ||  |||
cgacacagacgcagt-ctttgtgtgctcacacacacgtgct
```

- why:
  - similar sequence
  - infer homology
  - infer function

sequence → structure → function

# Global vs. Local

S	P	Q	-	R	T	G	K	C	C	W	I	A	G	P	G	I	L	H	R	M	S	L
S	G	A	L	R	C	S	W	N	D	-	I	A	G	P	C	A	Q	H	-	M	S	A

**Global:** Needleman-Wunsch; similar length, highly similar sequences

**Local:** Smith-Waterman; finds region(s) of highest similarity and build outward

# BLAST


## Basic Local Alignment Search Tool

- **idea:** find high scoring local alignments between query sequence and target database
- **assumption:** true match alignments very likely to contain *within them* very high scoring matches
  - heuristics theme:* search quickly for homologous regions and then do slow/exact alignments


# Traditional alignment methods

## BLAST


- (1) For the query find the list of high scoring words of length  $w$ .



- (2) Compare the word list to the database and identify exact matches.



- (3) For each word match, extend alignment in both directions to find alignments that score greater than score threshold  $S$ .





# BLAST Steps

## 1. Seeding

Query word ( $W = 3$ )



Query: GSDFWQETRASFGCSLAALLNKCKT**PQG**QRLVNQWIKQPLMDKNRIEERLNLVEAFGCATSWPI

Neighborhood  
words

PQG	18
PEG	15
PRG	14
PKG	14
PNG	13
PDG	13
PHG	13
PMG	13
PSG	13
PQA	12
PQN	12

Neighborhood  
score threshold ( $T = 13$ )

...  
Hit

Query: SLAALLNKCKT**PQG**QRLVNQWIKQPLMDKNRIEERLNLVEA

Subject: TLASVLDCTVT**PMG**SRMLKRWLHMPVRDTRVLLERQQTIGA

Determine the locations of all common “words” between the query and the database (“word hits”).

(protein  $W = 3$ , DNA  $W = 11$ )

# BLAST Steps

## 2. Extension

use dynamic programming to extend hits until  
the score drops a value of  $X$  – *expensive!!* --  
*90% of time*

ABCDEFGHIJKLMNOPQRSTUVWXYZ

| | | | | | | | | |

ABCDEFGHIJKLMNOPQRSTUVWXYZ

12345654567898765654

00000012100001234345

Match = 1

Mismatch = -1

**X = 5**

-> Score

-> Drop off score

# Scoring Matrices


A simple matrix for DNA

	C	T	A	G
C	1	-1	-1	-1
T	-1	1	-1	-1
A	-1	-1	1	-1
G	-1	-1	-1	1

ATGGCCATG  
 | | | | |  
 A-CGCCTCG

Score = 1

A more sophisticated matrix for DNA



	C	T	A	G
C	2	1	-1	-1
T	1	2	-1	-1
A	-1	-1	2	1
G	-1	-1	1	2

ATGGCCATG  
 | : | | | |  
 AC-GCCTCG

Score = 8

# Alignment Overview

- **Reads:** short DNA sequences usually up to 100-200 base pairs produced by a sequencing machine
- **Reference:** Genome sequence of organism of interest
- **Aligner:** Short-read aligner (BWA, bowtie, SOAP, MAQ etc.)
- **Distances:**

- ❑ **Hamming Distance:**

- The hamming distance is defined only for **strings of the same length**. For two strings, it is the number of places in which the two string differ.

- ❑ **Edit distance:**

- The edit distance between two strings is the minimum number of insertions, deletions and substitutions needed to transform the first string into the second one.

# Overview

MICHAEL  
MICHELE

A = E

E = L

L = E

Hamming distance = 3

KOBY  
BOBBY

K=B

Y=B

Y=inserted


Edit distance = 3

# The Burroughs Wheeler Alignment BWA tools

- BWA is used to map low divergent sequence reads to a reference genome
- It assumes that your 'reads' are from the genome you are aligning to
- There are three algorithms; the most common being BWA-MEM
- BWA outputs alignments in a SAM format
- Downstream you can use SAM tools or GATK to call variants or whatever

# Alignment methods for NGS

## Burrows-Wheeler transformation



# *Tools for NGS aligning*

## BWA

“There is no such thing (yet) as an automated gearshift in short read mapping. It is all like stick-shift driving in San Francisco. In other words running this tool with default parameters will probably not give you meaningful results. A way to deal with this is to understand the parameters by carefully reading the documentation and experimenting. Fortunately, Galaxy makes experimenting easy.”

- Galaxy Team

## Parameters to test





# Popular methods for alignment

## **BWA**

### **BWA-aln:**

- ☐ Short reads up to 200 bp with errors <5%
- ☐ Gapped alignment
- ☐ global alignment
- ☐ Can do paired-end
- ☐ Report ambiguous hits


### **BWA-SW:**

- ☐ Can align longer reads (upto 1mbp)
- ☐ Local alignment
- ☐ The paired-end mode only works for Illumina short-insert libraries.

**BWA-MEM is the latest BWA algorithm – much faster**

**BWA outputs the final alignment in the SAM (Sequence Alignment/Map) format**

# Mapping Quality



- Reads can occur more than once in the reference genome
- One can restrict the analysis to exclude the reads which occur more than  $n$  times
- As  $n$  gets larger, one gets more data, but also more noise in the data

# Mapping Quality

**Understanding mapping qualities:**

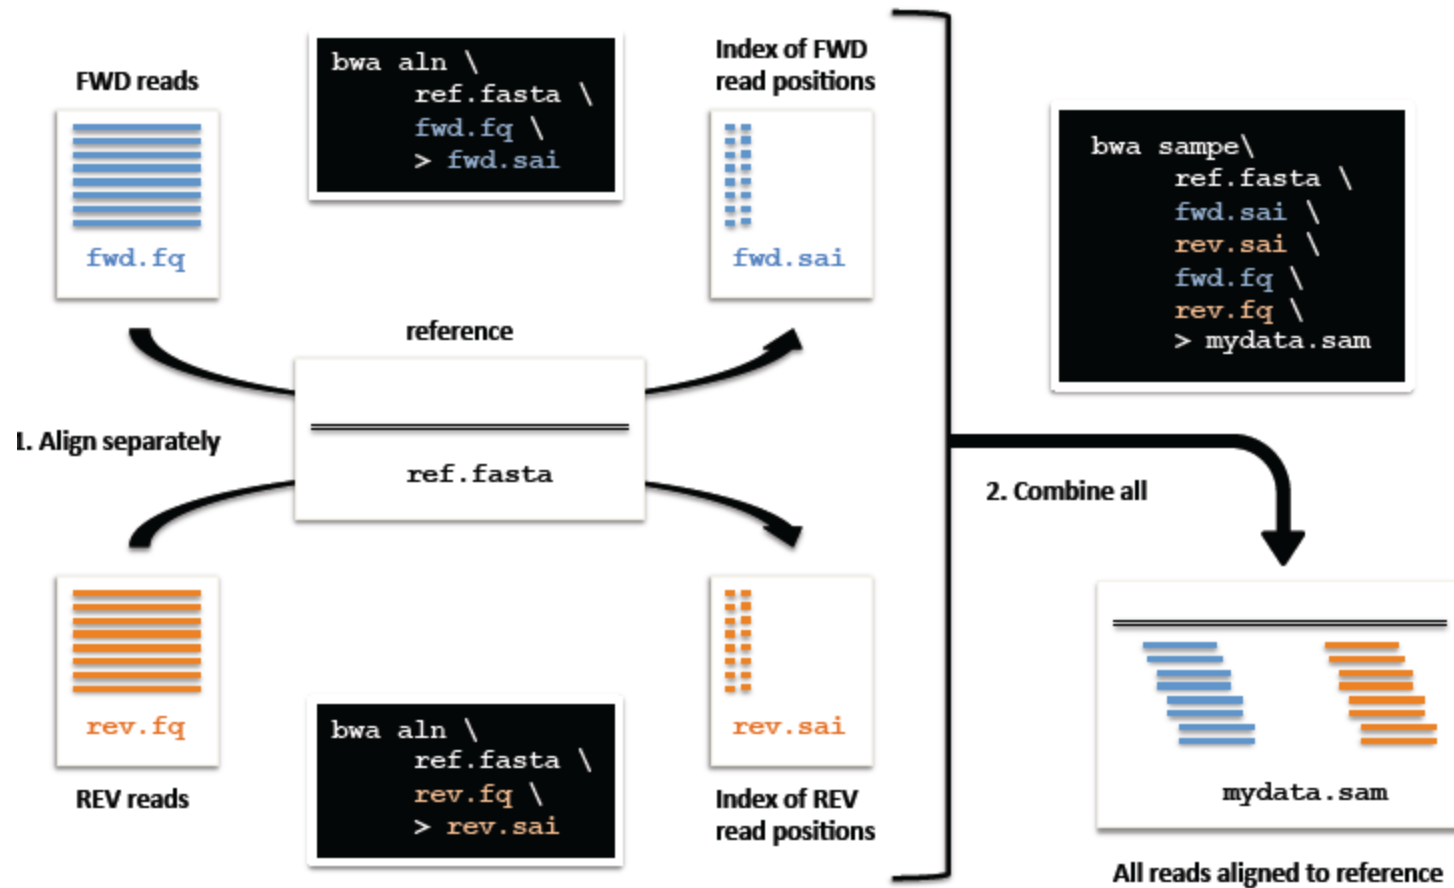
**Mapping quality calculation consider all these factors:**

- ☐ **Reference genome repeat structure**
- ☐ **Base qualities of read**
- ☐ **Paired end or not**

**$Q_s = 30$  implies there is a 1 in 1000 probability that the read is incorrectly mapped.**

- ☐ **The overall base quality of the read is good**
- ☐ **The best alignment has fewer mismatches**
- ☐ **The read has not matched to many places in genome**

# Typical workflow using BWA to map paired-end data



# Prototypical IGV screenshot representing aligned NGS reads

Non-reference bases are colored;  
reference bases are grey


Clean C/T  
heterozygote

Depth of coverage

First and second read from  
the same fragment

Individual reads aligned  
to the genome

Reference genome



# SAM format

**SAM stands for Sequence Alignment/Map format**

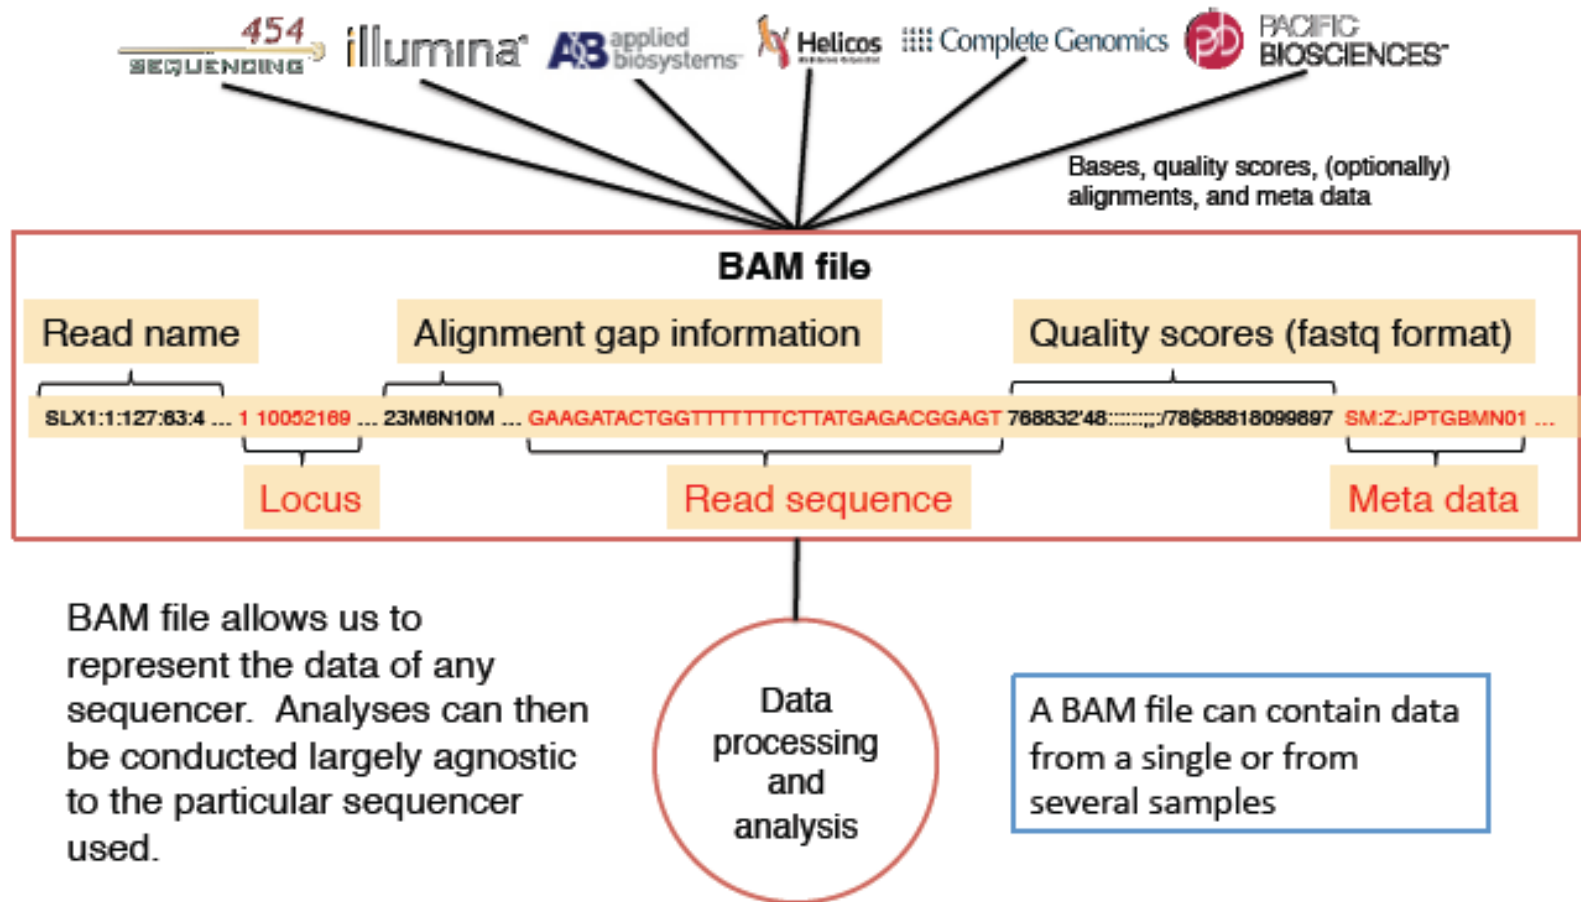
- ☐ **TAB-delimited text format consisting of optional header section and an alignment section**
- ☐ **Header lines start with “@” symbol; alignment lines don’t**
- ☐ **Each alignment line has 11 mandatory fields for essential alignment information**
- ☐ **Variable number of optional fields for flexible or aligner specific information.**

**It’s compact version is BAM format (Binary alignment MAP)**

# SAM format

- ☐ Alignment file SAM is converted to BAM format for efficient storage and access to alignment information
- ☐ BAM is also indexed to allow access to portions of information without loading the whole file
- ☐ BAM are re-ordered mostly by chromosomal coordinates
- ☐ Alignment paired or unpaired or different samples could be merged by samtools

# The BAM format stores aligned reads and is technology independent






# Cleaning up BAM Alignments

## Avoiding bad data

✖ = sequencing error propagated in duplicates



After marking duplicates, the GATK will only see :



... and thus be more likely to make the right call

# BAM headers: an essential part of a BAM file

@HD VN:1.0 GO:none SO:coordinate

@SQ SN:chrM LN:16571

@SQ SN:chr1 LN:247249719

@SQ SN:chr2 LN:242951149

[cut for clarity]

@SQ SN:chr9 LN:140273252

@SQ SN:chr10 LN:135374737

@SQ SN:chr11 LN:134452384

[cut for clarity]

@SQ SN:chr22 LN:49691432

@SQ SN:chrX LN:154913754

@SQ SN:chrY LN:57772954

@RG ID:20FUK.1 PL:illumina PU:20FUKAAXX100202.1 LB:Solexa-18483 SM:NA12878 CN:BI

@RG ID:20FUK.2 PL:illumina PU:20FUKAAXX100202.2 LB:Solexa-18484 SM:NA12878 CN:BI

@RG ID:20FUK.3 PL:illumina PU:20FUKAAXX100202.3 LB:Solexa-18483 SM:NA12878 CN:BI

@RG ID:20FUK.4 PL:illumina PU:20FUKAAXX100202.4 LB:Solexa-18484 SM:NA12878 CN:BI

@RG ID:20FUK.5 PL:illumina PU:20FUKAAXX100202.5 LB:Solexa-18483 SM:NA12878 CN:BI

@RG ID:20FUK.6 PL:illumina PU:20FUKAAXX100202.6 LB:Solexa-18484 SM:NA12878 CN:BI

@RG ID:20FUK.7 PL:illumina PU:20FUKAAXX100202.7 LB:Solexa-18483 SM:NA12878 CN:BI

@RG ID:20FUK.8 PL:illumina PU:20FUKAAXX100202.8 LB:Solexa-18484 SM:NA12878 CN:BI

@PG ID:BWA VN:0.5.7 CL:tk

@PG ID:GATK TableRecalibration VN:1.0.2864

20FUKAAXX100202.1:1:12730:189900 163 chrM 1 60 101M = 282 381

GATCACAGGTCTATCACCTATTAACCACTCACGGGAGCTCTCCATGCATTGGTA...[more bases]

?BA@A>BBBBACBBAC@BBBCBBCBC@BC@CAC@:BBBCBACACBABCBCB...[more quals]

RG:Z:20FUK.1 NM:i:1 SM:i:37 AM:i:37 MD:Z:72G28 MQ:i:60 PG:Z:BWA UQ:i:33

Required: Standard header

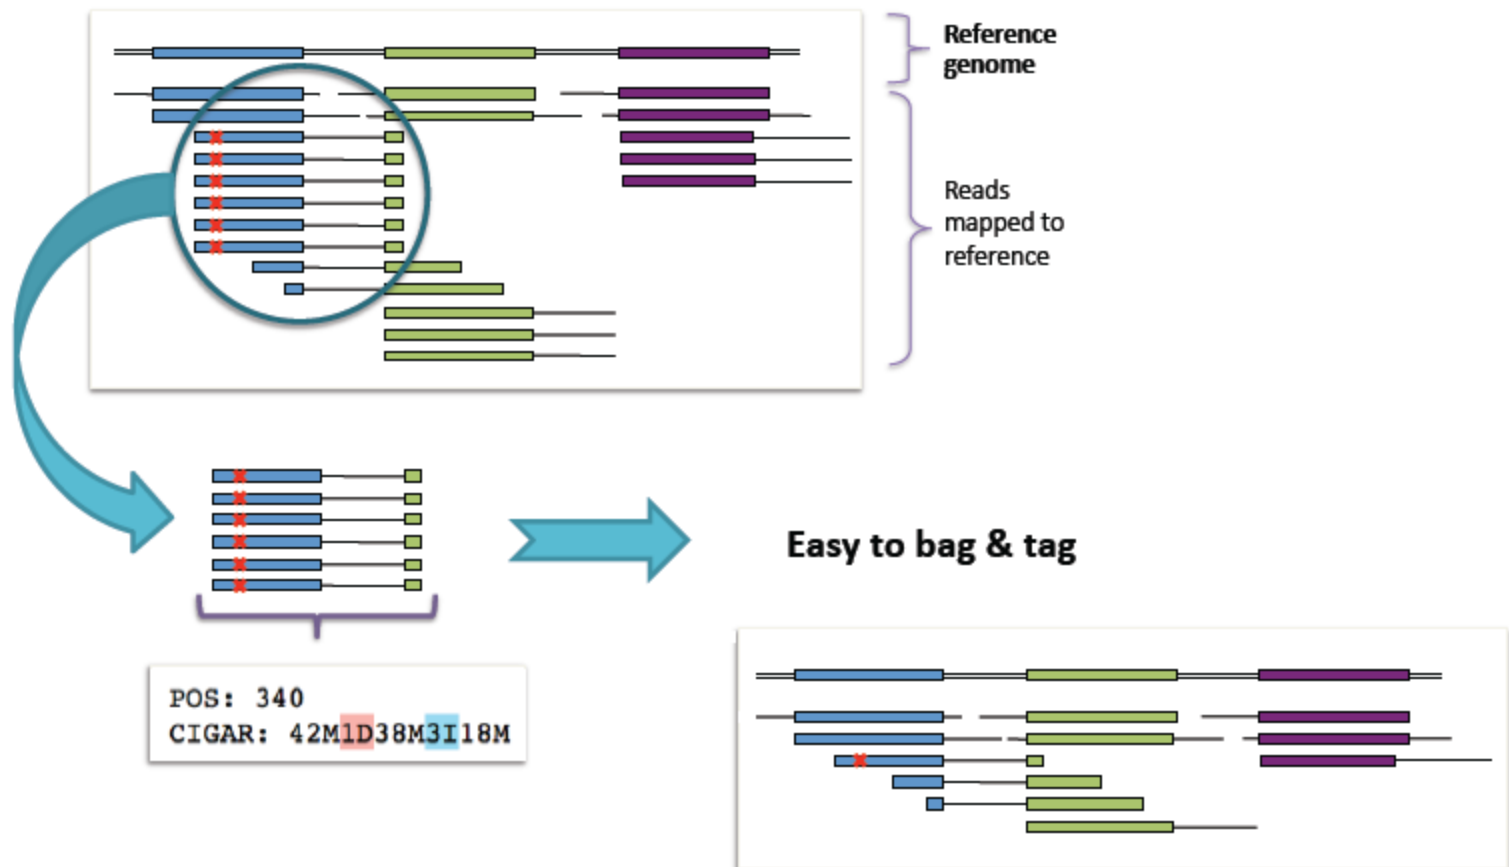
Essential: contigs of aligned reference sequence. Should be in karyotypic order.

Essential: read groups. Carries platform (PL), library (LB), and sample (SM) information. Each read is associated with a read group

Useful: Data processing tools applied to the reads

Official specification in <http://samtools.sourceforge.net/SAM1.pdf>

# Identifying duplicates



# VCF Files store variant information

```
##fileformat=VCFv4.1
##reference=1000GenomesPilot-NCBI36
##INFO=<ID=DP,Number=1,Type=Integer,Description="Total Depth">
##INFO=<ID=AF,Number=A,Type=Float,Description="Allele Frequency">
##INFO=<ID=DB,Number=0,Type=Flag,Description="dbSNP membership, build 129">
##FILTER=<ID=s50,Description="Less than 50% of samples have data">
##FORMAT=<ID=GT,Number=1,Type=String,Description="Genotype">
##FORMAT=<ID=GQ,Number=1,Type=Integer,Description="Genotype Quality">
##FORMAT=<ID=DP,Number=1,Type=Integer,Description="Read Depth">
#CHROM POS ID REF ALT QUAL FILTER INFO
FORMAT NA000001 NA000002 NA000003

20 14370 rs6054257 G A 29 PASS DP=14;AF=0.5;DB
GT:GQ:DP 0|0:48:1 1|0:48:8 1/1:43:5
20 1110696 rs6040355 A G,T 67 PASS DP=10;AF=0.333,0.667;DB
GT:GQ:DP 1|2:21:6 2|1:2:0 2/2:35:4
20 1230237 . T . 47 PASS DP=13
GT:GQ:DP 0|0:54:7 0|0:48:4 0/0:61:2
20 1234567 microsat1 GTCT G,GTACT 50 PASS DP=9
GT:GQ:DP 0/1:35:4 0/2:17:2 1/1:40:3
```

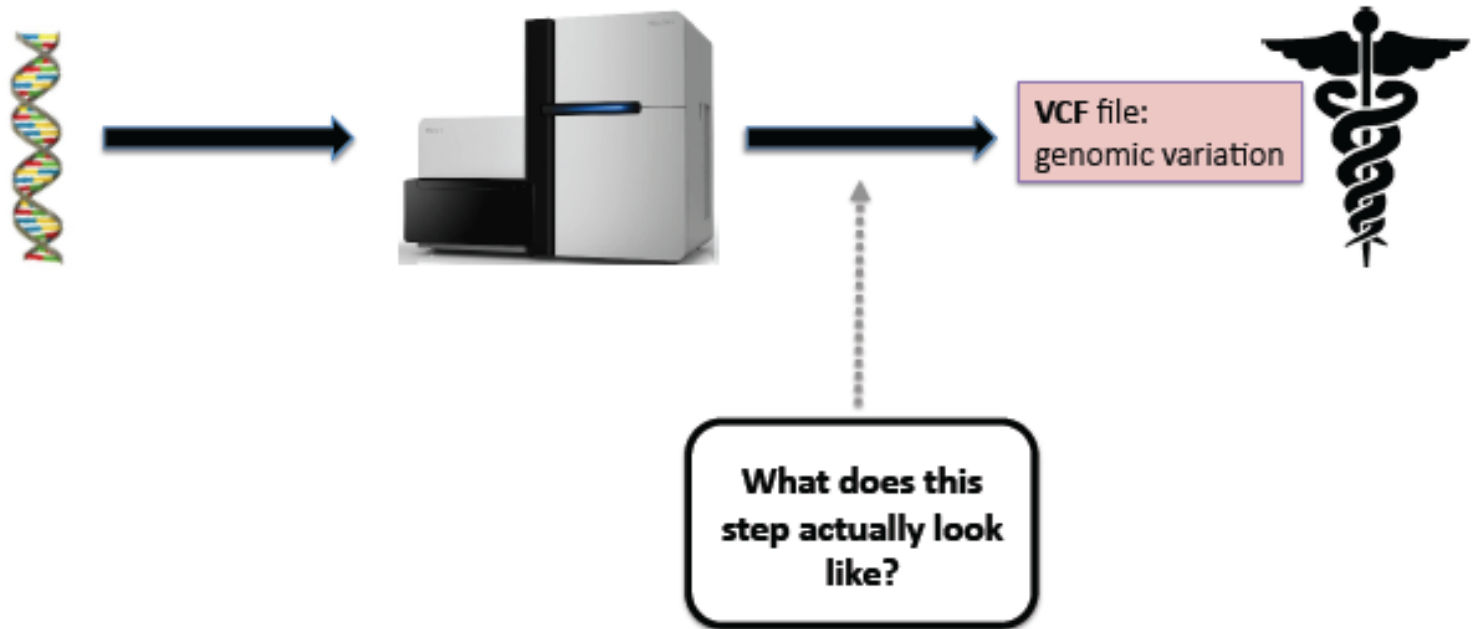
Header

Variant  
records


Official specification in

[www.1000genomes.org/wiki/Analysis/Variant Call Format/vcf-variant-call-format-version-41](http://www.1000genomes.org/wiki/Analysis/Variant%20Call%20Format/vcf-variant-call-format-version-41)

Is processing/analysis of NGS data really that easy?




It's going to involve dealing with messy situations like this:



How can we tell which mismatches represent real mutations and which are just noise?

# If you look at a lot of data, you're bound to find something

**Number of people who drowned by falling into a pool**  
correlates with  
**Films Nicolas Cage appeared in**



<http://www.tylervigen.com/spurious-correlations>

